

Introduction to C++: Part 4

Tutorial Outline: Part 3

- Virtual functions and inheritance
- Look at a real and useful C++ class

Square

- Let's make a subclass of Rectangle called Square.
- Open the NetBeans project *Shapes*
- This has the Rectangle class from Part 2 implemented.
- Add a class named *Square*.
- Make it inherit from Rectangle.

Square.h

```
#ifndef SQUARE_H
#define SQUARE_H

#include "Rectangle.h"

class Square : public Rectangle
{
    public:
        Square();
        virtual ~Square();

    protected:

    private:
};

#endif // SQUARE_H
```

Square.cpp

```
#include "Square.h"

Square::Square()
{}

Square::~~Square()
{}

```

- Note that subclasses are free to add any number of new methods or members, they are not limited to those in the superclass.

- Class Square inherits from class Rectangle

A new Square constructor is needed.

- A square is, of course, just a rectangle with equal length and width.
- The area can be calculated the same way as a rectangle.
- Our Square class therefore needs just one value to initialize it and it can re-use the Rectangle.Area() method for its area.
- Go ahead and try it:
 - Add an argument to the default constructor in Square.h
 - Update the constructor in Square.cpp to do...?
 - Remember Square can access the public members and methods in its superclass



Solution 1

```
#ifndef SQUARE_H
#define SQUARE_H

#include "Rectangle.h"

class Square : public Rectangle
{
    public:
        Square(float width);
        virtual ~Square();

    protected:

    private:
};

#endif // SQUARE_H
```


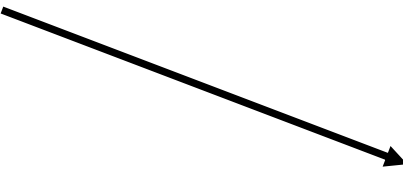
```
#include "Square.h"

Square::Square(float length) :
    m_width (length), m_length (length)
{

}
```

- Square can access the public members in its superclass.
- Its constructor can then just assign the length of the side to the Rectangle m_width and m_length.
- This is unsatisfying – while there is nothing *wrong* with this it's not the OOP way to do things.
- Why re-code the perfectly good constructor in Rectangle?

The delegating constructor

- C++11 added a new constructor type called the delegating constructor.
- Using member initialization lists you can call one constructor from another. 
- Even better: with member initialization lists C++ can call superclass constructors! 

```
class class_c {  
public:  
    int max;  
    int min;  
    int middle;  
  
    class_c(int my_max) {  
        max = my_max > 0 ? my_max : 10;  
    }  
    class_c(int my_max, int my_min) : class_c(my_max) {  
        min = my_min > 0 && my_min < max ? my_min : 1;  
    }  
    class_c(int my_max, int my_min, int my_middle) :  
        class_c(my_max, my_min) {  
        middle = my_middle < max &&  
            my_middle > min ? my_middle : 5;  
    }  
};
```

Reference:

<https://msdn.microsoft.com/en-us/library/dn387583.aspx>

```
Square::Square(float length) :  
    Rectangle(length, length)  
{  
    // other code could go here.  
}
```

Solution 2

```
#ifndef SQUARE_H
#define SQUARE_H

#include "Rectangle.h"

class Square : public Rectangle
{
    public:
        Square(float width);
        virtual ~Square();

    protected:

    private:
};

#endif // SQUARE_H
```

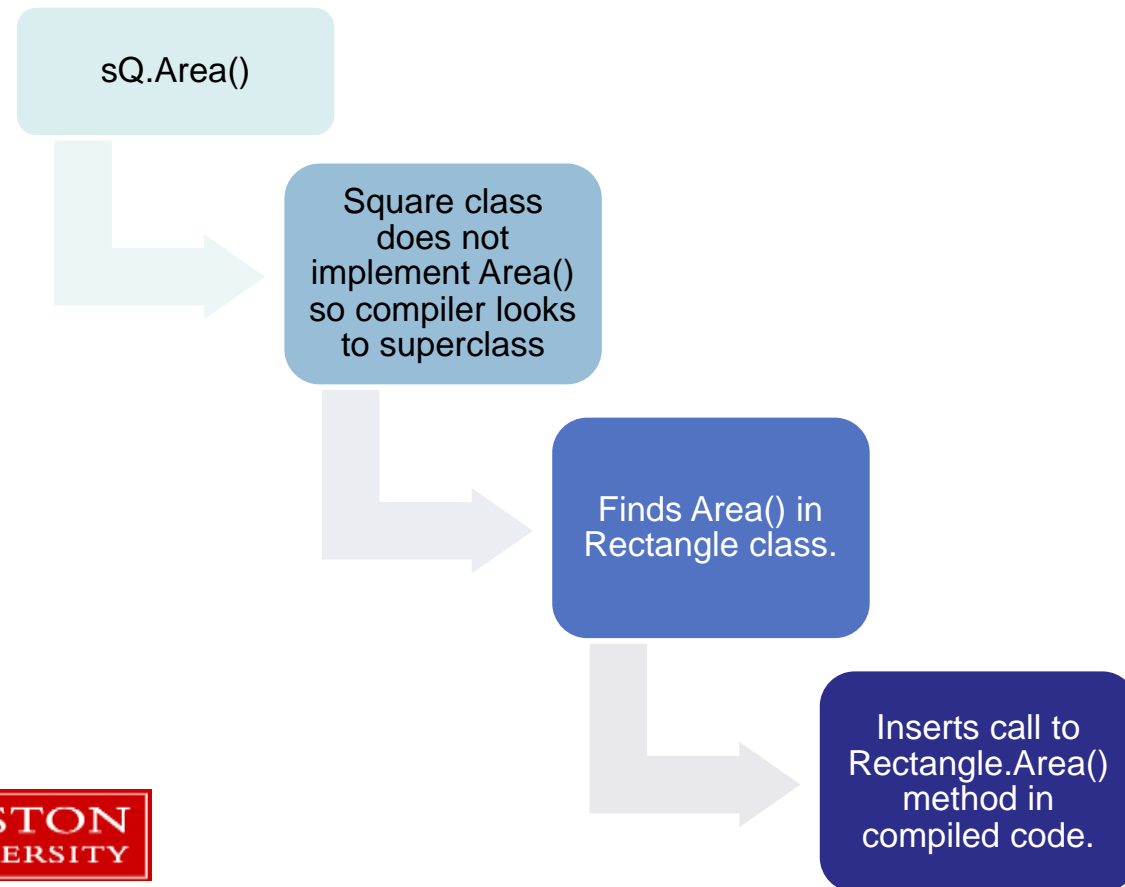
```
#include "Square.h"

Square::Square(float length) :
    Rectangle(length, length) {}
```

- Square can directly call its superclass constructor and let the Rectangle constructor make the assignment to m_width and m_length.
- This saves typing, time, and **reduces the chance of adding bugs to your code.**
 - The more complex your code, the more compelling this statement is.
- Code re-use is one of the prime reasons to use OOP.

Trying it out in main()

- What happens behind the scenes when this is compiled....



```
#include <iostream>

using namespace std;

#include "Square.h"

int main()
{
    Square sQ(4) ;

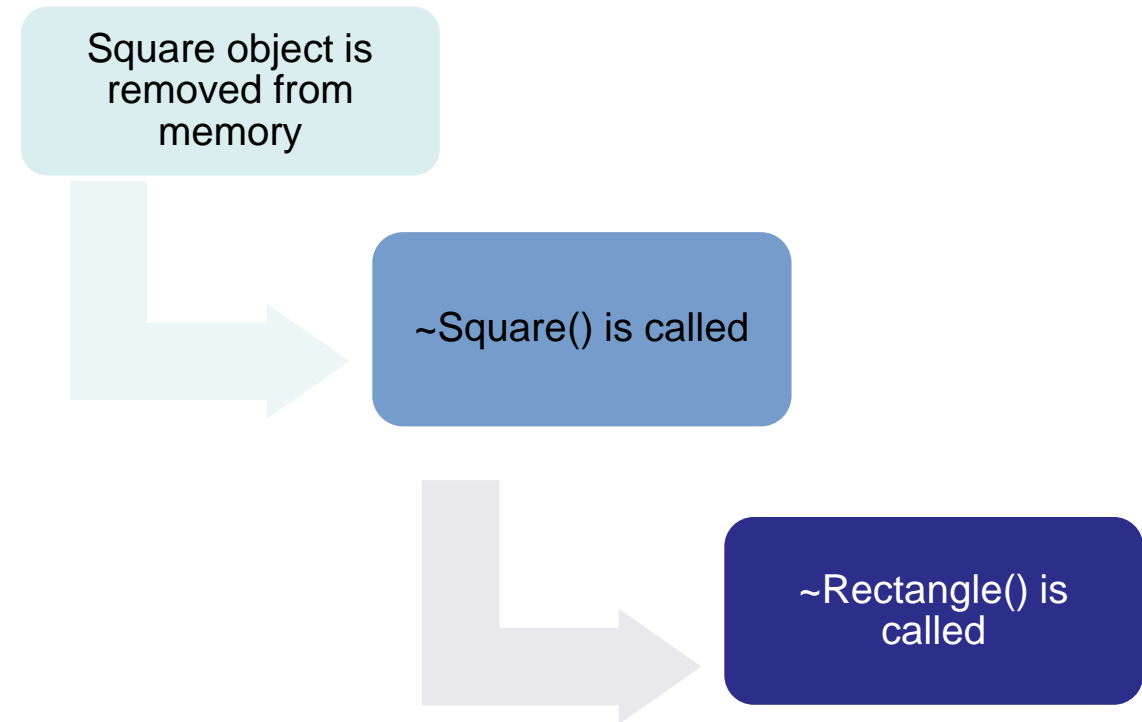
    // Uses the Rectangle Area() method!
    cout << sQ.Area() << endl ;

    return 0;
}
```



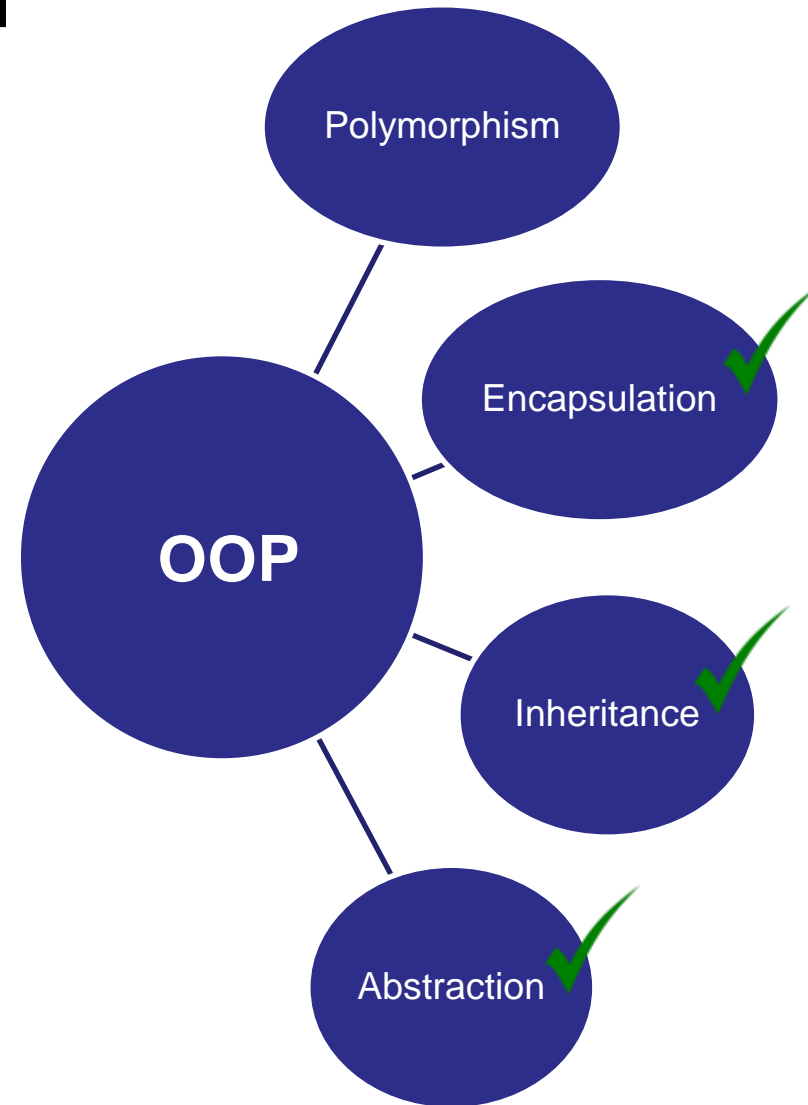
More on Destructors

- When a subclass object is removed from memory, its destructor is called as it is for any object.
- Its superclass destructor is then *also* called .
- Each subclass should only clean up its own problems and let superclasses clean up theirs.



The formal concepts in OOP

- Next up: Polymorphism



Using subclasses

- A function that takes a superclass argument can *also* be called with a subclass as the argument.
- The reverse is **not** true – a function expecting a subclass argument cannot accept its superclass.
- Copy the code to the right and add it to your main.cpp file.

```
void PrintArea(Rectangle &rT) {  
    cout << rT.Area() << endl ;  
}  
  
int main() {  
    Rectangle rT(1.0,2.0) ;  
    Square sQ(3.0) ;  
    PrintArea(rT) ;  
    PrintArea(sQ) ;  
}
```

The PrintArea function can accept the Square object sQ because Square is a subclass of Rectangle.



Overriding Methods

- Sometimes a subclass needs to have the same interface to a method as a superclass but with different functionality.
- This is achieved by *overriding* a method.
- Overriding a method is simple: just re-implement the method with the same name and arguments in the subclass.

```
class Super {
public:
    void PrintNum() {
        cout << 1 << endl ;
    }
};

class Sub : public Super {
public:
    // Override
    void PrintNum() {
        cout << 2 << endl ;
    }
};

Super sP ;
sP.PrintNum() ; // Prints 1
Sub sB ;
sB.PrintNum() ; // Prints 2
```



Overriding Methods

- Seems simple, right?

```
class Super {  
public:  
    void PrintNum() {  
        cout << 1 << endl ;  
    }  
};  
  
class Sub : public Super {  
public:  
    // Override  
    void PrintNum() {  
        cout << 2 << endl ;  
    }  
};  
  
Super sP ;  
sP.PrintNum() ; // Prints 1  
Sub sB ;  
sB.PrintNum() ; // Prints 2
```

How about in a function call...

- Using a single function to operate on different types is *polymorphism*.
- Given the class definitions, what is happening in this function call?

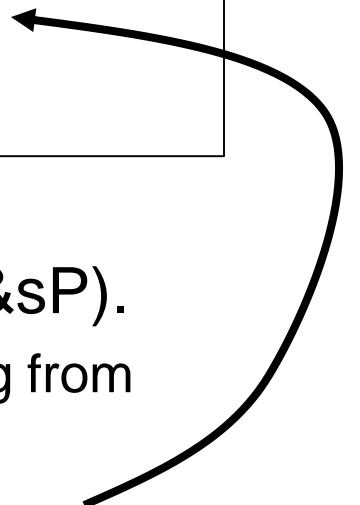
“C++ is an insult to the human brain”
– Niklaus Wirth (designer of Pascal)

```
class Super {  
public:  
    void PrintNum() {  
        cout << 1 << endl ;  
    }  
};  
  
class Sub : public Super {  
public:  
    // Override  
    void PrintNum() {  
        cout << 2 << endl ;  
    }  
};
```

```
void FuncRef(Super &sP) {  
    sP.PrintNum() ;  
}  
  
Super sP ;  
Func(sP) ; // Prints 1  
Sub sB ;  
Func(sB) ; // Hey!! Prints 1!!
```

Type casting

```
void FuncRef (Super &sP) {  
    sP.PrintNum() ;  
}
```



- The Func function passes the argument as a *reference* (Super &sP).
 - What's happening here is *dynamic type casting*, the process of converting from one type to another at runtime.
 - Same mechanism as the *dynamic_cast<type>()* function
- The incoming object is treated as though it were a superclass object in the function.
- When methods are overridden and called there are two points where the proper version of the method can be identified: either at compile time or at runtime.

Virtual methods

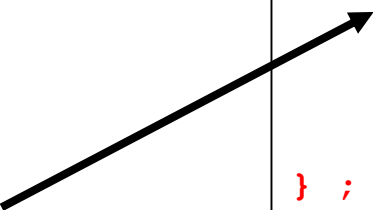
- When a method is labeled as virtual and overridden the compiler will generate code that will check the type of an object at **runtime** when the method is called.
- The type check will then result in the expected version of the method being called.
- When overriding a virtual method in a subclass, it's a good idea to label the method as virtual in the subclass as well.
 - ...just in case this gets subclassed again!

```
class SuperVirtual
{
public:
    virtual void PrintNum()
    {
        cout << 1 << endl ;
    }
};

class SubVirtual : public SuperVirtual
{
public:
    // Override
    virtual void PrintNum()
    {
        cout << 2 << endl ;
    }
};

void Func(SuperVirtual &sP)
{
    sP.PrintNum() ;
}

SuperVirtual sP ;
Func(sP) ; // Prints 1
SubVirtual sB ;
Func(sB) ; // Prints 2!!
```




Early (static) vs. Late (dynamic) binding

- Leaving out the virtual keyword on a method that is overridden results in the compiler deciding *at compile time* which version (subclass or superclass) of the method to call.
- This is called early or static *binding*.
- At compile time, a function that takes a superclass argument will only call the **non-virtual** superclass method under early binding.
- Making a method virtual adds code behind the scenes (that you, the programmer, never interact with directly)
 - Lookups in a hidden table, called the *vtable*, are done to figure out what version of the virtual method should be run.
- This is called late or dynamic binding.
- There is a small performance penalty for late binding due to the vtable lookup.
- **This only applies when an object is referred to by a reference or pointer.**

Behind the scenes – vptr and vtable

- C++ classes have a hidden pointer (vptr) generated that points to a table of virtual methods associated with a class (vtable).
- When a virtual class method (base class or its subclasses) is called by reference (or pointer) *when the program is running* the following happens:
 - The object's **class** vptr is followed to its **class** vtable
 - The virtual method is looked up in the vtable and is then called.
 - One vptr and one vtable per class so minimal memory overhead
 - If a method override is **non-virtual** it won't be in the vtable and it is selected at **compile time**.



Let's run this through the debugger

- Open the project `Virtual_Method_Calls`.
- Everything here is implemented in one big `main.cpp`
- Place a breakpoint at the first line in `main()` and in the two implementations of `Func()`



When to make methods virtual

- If a method will be (or might be) overridden in a subclass, make it virtual
 - There is a *minuscule* performance penalty. Will that even matter to you?
 - i.e. Have you profiled and tested your code to show that virtual method calls are a performance issue?
 - When is this true?
 - Almost always! Who knows how your code will be used in the future?
- Constructors are **never** virtual in C++.
- Destructors in a base class should always be virtual.
 - Also – if any method in a class is virtual, make the destructor virtual
 - These are important when dealing with objects via reference and it avoids some subtleties when manually allocating memory.

Why all this complexity?

```
void FuncEarly(SuperVirtual &sP)
{
    sP.PrintNum();
}
```

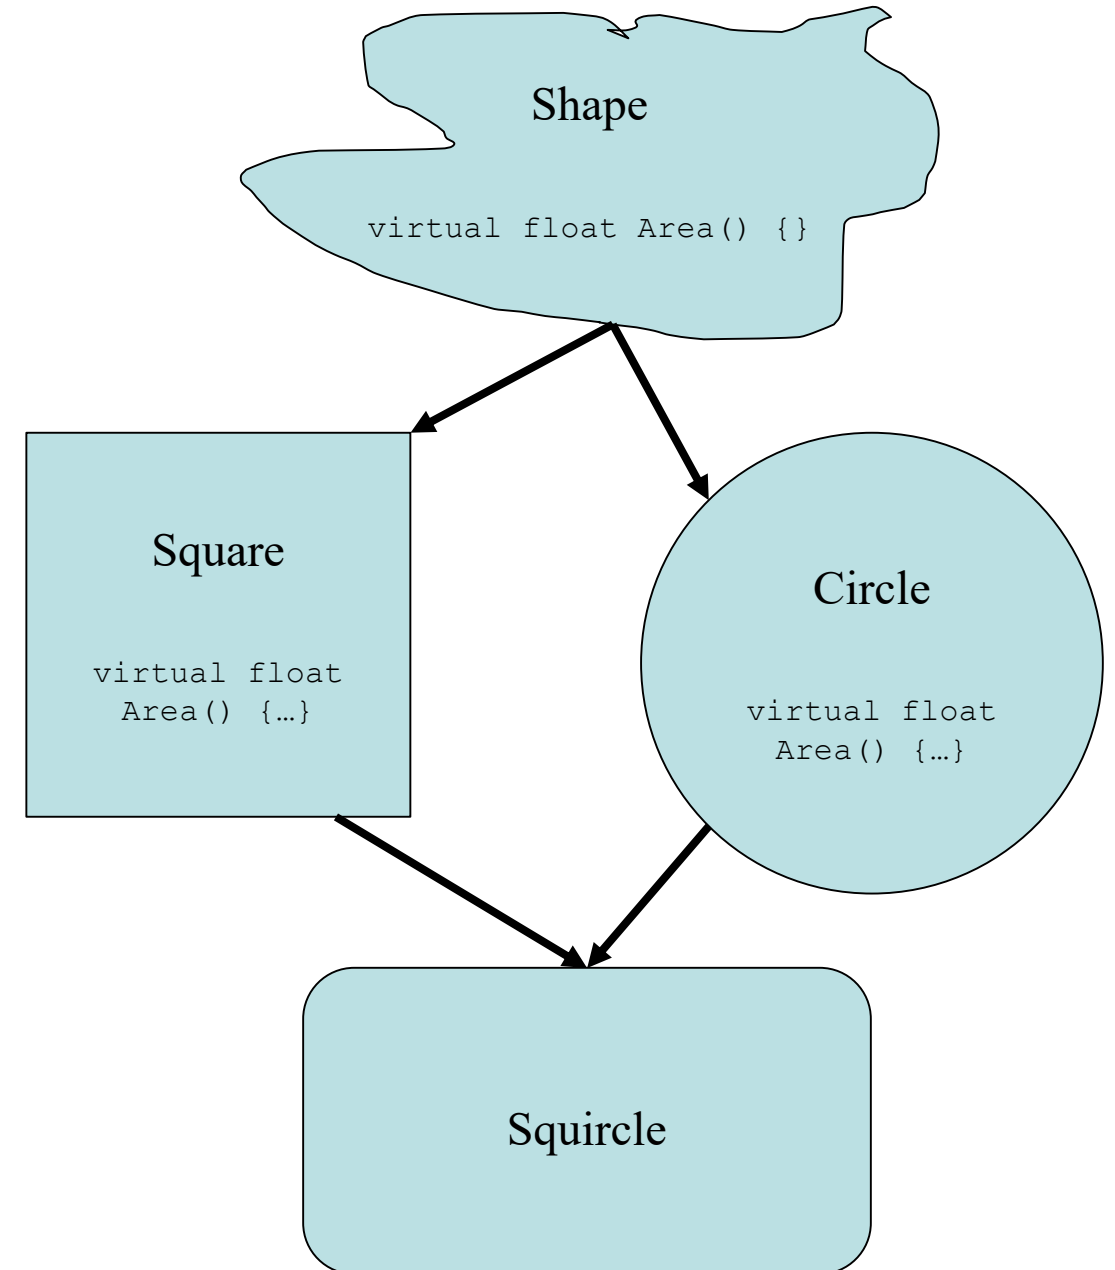
- Called by **reference** – late binding to PrintNum()

```
void FuncLate(SuperVirtual sP)
{
    sP.PrintNum();
}
```

- Called by **value** – early binding to PrintNum even though it's virtual!


- Late binding allows for code libraries to be updated for new functionality. As methods are identified at runtime the executable does not need to be updated.
- This is done all the time! Your C++ code may be, for example, a plugin to an existing simulation code.
- Greater flexibility when dealing with multiple subclasses of a superclass.
- Most of the time this is the behavior you are looking for when building class hierarchies.

- Remember the Deadly Diamond of Death? Let's explain.
- Look at the class hierarchy on the right.
 - Square and Circle inherit from Shape
 - Squirrel inherits from both Square and Circle
 - Syntax:
class Squirrel : public Square, public Circle
- The Shape class implements an empty Area() method. The Square and Circle classes override it. Squirrel does not.
- Under late binding, which version of Area is accessed from Squirrel?
Square.Area() or Circle.Area()?



Interfaces

- Interfaces are a way to have your classes share behavior without them sharing actual code.
- Gives much of the benefit of multiple inheritance without the complexity and pitfalls



- Example: for debugging you want each class to have a Log() method that writes some info to a file.
 - Implement with an interface.

Interfaces

- An interface class in C++ is called a pure virtual class.
- It contains virtual methods only with a special syntax. Instead of {} the function is set to 0.
 - Any subclass needs to implement the methods!
- Modified Square.h shown.
- What happens when this is compiled?

```
(...error...)
include/square.h:10:7: note:   because the following virtual
functions are pure within 'Square':
  class Square : public Rectangle, Log
    ^
include/square.h:7:18: note:   virtual void Log::LogInfo()
    virtual void LogInfo()=0 ;
```

- Once the LogInfo() is uncommented it will compile.

```
#ifndef SQUARE_H
#define SQUARE_H

#include "rectangle.h"


class Log {
    virtual void LogInfo()=0 ;
};

class Square : public Rectangle, Log
{
    public:
        Square(float length) ;
        virtual ~Square() ;
        // virtual void LogInfo() {}
    protected:
        private:
};

#endif // SQUARE_H
```

Putting it all together

- Now let's revisit our Shapes project.
- Open the “**Shapes with Circle**” project.
 - This has a Shape base class with a Rectangle and a Square
- Add a Circle class to the class hierarchy in a sensible fashion.



- Hint: Think first, code second.



New pure virtual Shape class

- Slight bit of trickery:
 - An empty constructor is defined in shape.h
 - No need to have an extra shape.cpp file if these functions do nothing!
- Q: How much code can be in the header file?
- A: Most of it with some exceptions.
 - .h files are not compiled into .o files so a header with a lot of code gets re-compiled every time it's referenced in a source file.

```
#ifndef SHAPE_H
#define SHAPE_H

class Shape
{
    public:
        Shape() {}
        virtual ~Shape() {}

        virtual float Area()=0 ;
    protected:

    private:
};


#endif // SHAPE_H
```

Give it a try

- Add inheritance from Shape to the Rectangle class
- Add a Circle class, inheriting from wherever you like.
- Implement Area() for the Circle
- If you just want to see a solution, open the project “Shapes with Circle solved”


A Potential Solution

- A Circle has one dimension (radius), like a Square.
 - Would only need to override the Area() method
- But...
 - Would be storing the radius in the members m_width and m_length. This is not a very obvious to someone else who reads your code.
- Maybe:
 - Change m_width and m_length names to m_dim_1 and m_dim_2?
 - Just makes everything more muddled!



A Better Solution

- Inherit separately from the Shape base class
 - Seems logical, to most people a circle is not a specialized form of rectangle...
- Add a member `m_radius` to store the radius.
- Implement the `Area()` method
- Makes more sense!
- Easy to extend to add an Oval class, etc.



New Circle class

- Also inherits from Shape
- Adds a constant value for π
 - Constant values can be defined right in the header file.
 - If you accidentally try to change the value of PI the compiler will throw an error.

```
#ifndef CIRCLE_H
#define CIRCLE_H

#include "shape.h"

class Circle : public Shape
{
    public:
        Circle();
        Circle(float radius) ;
        virtual ~Circle();

        virtual float Area() ;

        const float PI = 3.14;
        float m_radius ;

    protected:

    private:
};

#endif // CIRCLE_H
```

- circle.cpp
- Questions?

```
#include "circle.h"

Circle::Circle()
{
    //ctor
}

Circle::~~Circle()
{
    //dtor
}

// Use a member initialization list.
Circle::Circle(float radius) : m_radius{radius}
{}

float Circle::Area()
{
    // Quiz: what happens if this line is
    // uncommented and then compiled:
    //PI=3.14159 ;
    return m_radius * m_radius * PI ;
}
```


Quiz time!

- What happens behind the scenes when the function PrintArea is called?
- How about if PrintArea's argument was instead:

```
void PrintArea(Shape shape)
```

```
void PrintArea(Shape &shape) {  
    cout << "Area: " << shape.Area() << endl ;  
}  
  
int main()  
{  
    Square sQ(4) ;  
    Circle circ(3.5) ;  
    Rectangle rT(21,2) ;  
  
    // Print everything  
    PrintArea(sQ) ;  
    PrintArea(rT) ;  
    PrintArea(circ) ;  
    return 0;  
}
```

Quick mention...

- Aside from overriding functions it is also possible to override operators in C++.

- As seen in the C++ string. The + operator concatenates strings:

```
string str = "ABC" ;  
str = str + "DEF" ;  
// str is now "ABCDEF"
```

- It's possible to override +,-,=,<,>, brackets, parentheses, etc.

- Syntax:

```
MyClass operator*(const MyClass& mC) {...}
```

- Recommendation:

- Generally speaking, avoid this. This is an easy way to generate very confusing code.
 - A well-named function will almost always be easier to understand than an operator.

- An exceptions is the assignment operator: operator=

Summary

- C++ classes can be created in hierarchies via inheritance, a core concept in OOP.
- Classes that inherit from others can make use of the superclass' public and protected members and methods
 - You write less code!
- Virtual methods should be used whenever methods will be overridden in subclasses.
- Avoid multiple inheritance, use interfaces instead.
- Subclasses can override a superclass method for their own purposes and can still explicitly call the superclass method.
- Abstraction means hiding details when they don't need to be accessed by external code.
 - Reduces the chances for bugs.
- While there is a lot of complexity here – in terms of concepts, syntax, and application – keep in mind that OOP is a highly successful way of building programs!

A high quality random number generator

- The motivation for this code can be found on the RCS website: http://rcs.bu.edu/examples/random_numbers/
- The RNG implemented here is the xoroshiro128+ algorithm. The inventors published a C implementation on their website:


<http://xoshiro.di.unimi.it/>

Some OOP Guidelines

- Here are some guidelines for putting together a program using OOP to keep in mind while getting up and running with C++.
- Keep your classes simple and single purpose.
- Logically organize your classes to re-use code via inheritance.
- Use interfaces in place of multiple inheritance
- Keep your methods short
 - Many descriptive methods that do little things is easier to debug and understand.
- Follow the KISS principle:
 - “Keep it simple stupid”
 - “Keep it simple, silly”
 - “Keep it short and sweet”
 - “Make Simple Tasks Simple!” – Bjarne Stroustrup
 - “Make everything as simple as possible, but not simpler” – Albert Einstein

Putting your classes together

- Effective use of OOP demands that the programmer think/plan/design first and code second.
- There is a large body of information on this topic:



- As this is an academic institution your code may:
 - Live on in your lab long after you have graduated
 - Be worked on by multiple researchers
 - Adapted to new problems you haven't considered
 - Be shared with collaborators
- For more structured environments (ex. a team of professional programmers) there exist concepts like SOLID:
 - [https://en.wikipedia.org/wiki/SOLID_\(object-oriented_design\)](https://en.wikipedia.org/wiki/SOLID_(object-oriented_design))
 - ...and there are many others.

Keep your classes simple

- Avoid “monster” classes that implement everything including the kitchen sink.
- Our Rectangle class just holds dimensions and calculates its area.
 - It cannot print out its area, send email, draw to the screen, etc.

- **Single responsibility principle:**

- Every class has responsibility for one piece of functionality in the program.
- https://en.wikipedia.org/wiki/Single_responsibility_principle
- Example:
 - An Image class holds image data and can read and write it from disk.
 - A second class, ImageFilter, has methods that manipulate Image objects and return new ones.

- **Resource Allocation Is Initialization (RAII):**

- A late 80’s concept, widely used in OOP.
- https://en.wikipedia.org/wiki/Resource_acquisition_is_initialization
- ALL Resources in a class are created in the constructor and released in the destructor.
 - Example: opening files, allocating memory, etc.
- If an object is created it is ready to use.

C++ Libraries

- There are a **LOT** of libraries available for C++ code.
 - Sourceforge alone has >7400
 - <https://sourceforge.net/directory/language:c++/os:windows/?q=library>
- Before jumping into writing your code, consider what you need and see if there are libraries available.
- Many libraries contain code developed by professionals or experts in a particular field.
- Consider what you are trying to accomplish in your research:
 - A) accomplishments in your field or
 - B) C++ programming?

C++ Compilers on the SCC

| Module name | Vendor | Compiler | Versions | C++11 support |
|-------------|-------------------------|----------|---------------|---------------|
| gnu | GNU | g++ | 4.4.7 - 7.2.0 | 4.9.2 & up |
| intel | Intel | icpp | 2016 – 2018 | 2017 & 2018 |
| pgi | Portland Group / Nvidia | pgc++ | 13.5 – 18.4 | 18.4 |
| llvm | LLVM | clang++ | 3.9 – 6.0 | All |

- There are 4 families of compilers on the SCC for C++.
 - To see versions use the *module avail* command, e.g. `module avail gnu`
- They have their strengths and weaknesses. For numeric code the intel and pgi compilers tend to produce the fastest code.
- For info on how to choose compiler optimizations for the SCC see the RCS website:
<http://www.bu.edu/tech/support/research/software-and-programming/programming/compilers/compiler-optimizations/>

Multithreading

- OpenMP
 - Open MP is a standard approach to writing multithreaded code to exploit multiple CPU cores with your program.
 - Fully supported in C++
 - See <http://www.openmp.org/> for details, or take an RCS tutorial on using it.
- Intel Thread Building Blocks
 - C++ specific library
 - Available on the SCC from Intel and is also open source.
 - Much more flexible and much more C++-ish than OpenMP
 - Offers high performance memory allocators for multithreaded code
 - Includes concurrent data types (vectors, etc.) that can automatically be shared amongst threads with no added effort for the programmer to control access to them.
 - If you want to use this and need help: help@scc.bu.edu



Math and Linear Algebra

- Eigen
 - http://eigen.tuxfamily.org/index.php?title=Main_Page
 - Available on the SCC.
 - “Eigen is a C++ template library for linear algebra: matrices, vectors, numerical solvers, and related algorithms.”
- Armadillo
 - <http://arma.sourceforge.net/>
 - Available on the SCC.
 - “Armadillo is a high quality linear algebra library (matrix maths) for the C++ language, aiming towards a good balance between speed and ease of use. Provides high-level syntax (API) deliberately similar to Matlab.”
 - Also see *matlab2cpp* (<https://github.com/jonathf/matlab2cpp>), a semi-automatic tool for converting Matlab code to C++ with Armadillo.
 - And also see *PyJet* (<https://github.com/wolfv/pyjet>), which converts Python and Numpy code to Armadillo/C++ code.
- OpenCV
 - <https://opencv.org>
 - A computer vision and image processing library, with excellent high-performance support for linear algebra, many algorithms, and GPU acceleration.